Lecture #11:

Convolutional Neural Networks

June, 8

Convolutional Neural Networks were inspired by real neurons

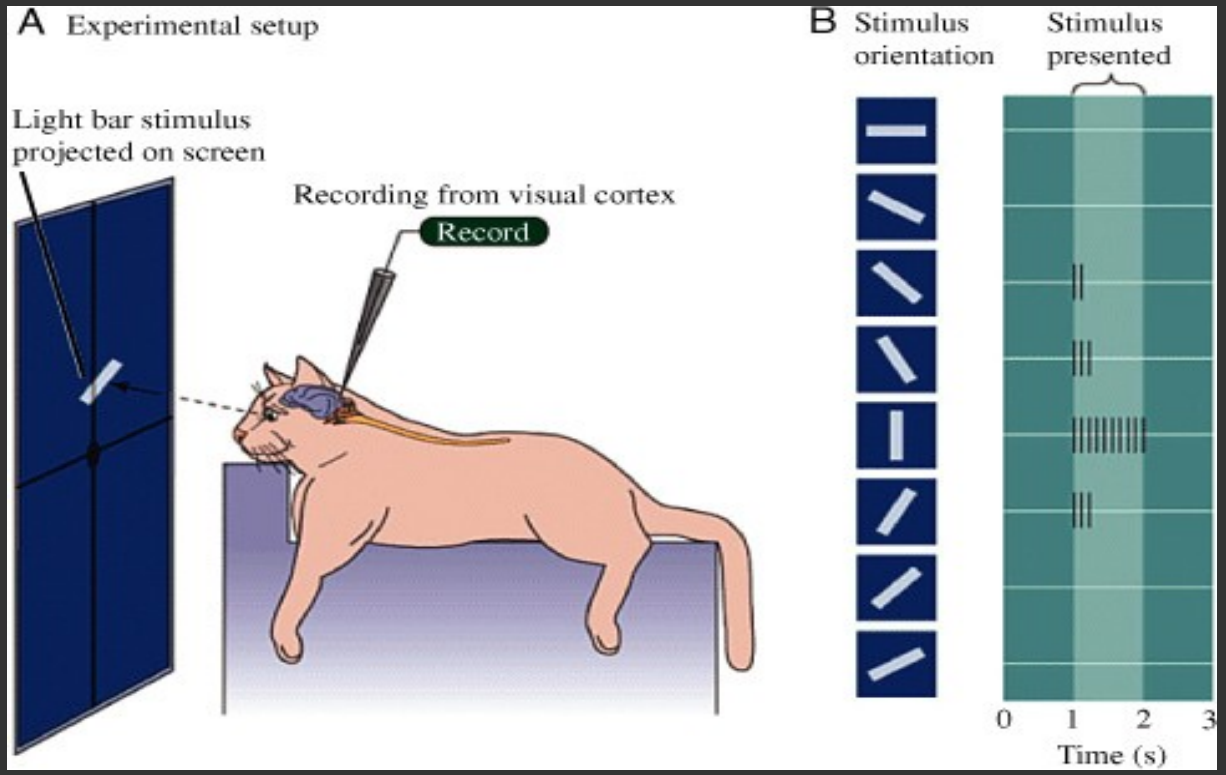

Hubel & Wiesel Cat Experiment, 1959

How images are stored

- Each image is represented by a matrix of pixels
- Each pixel is represented by 3 (or 4) channels
  3: Red, Green, Blue – a number from 0 to 255 (00 to ff)
  4: alpha channel, transparency

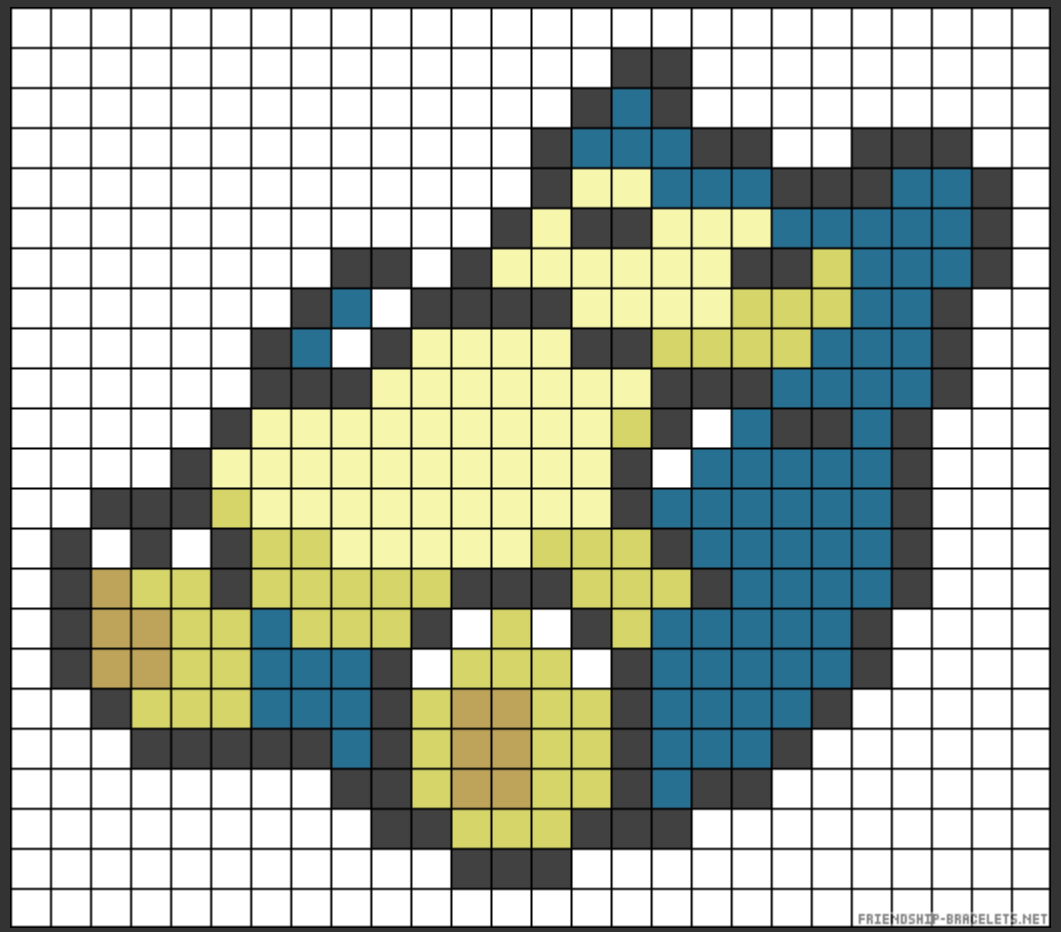

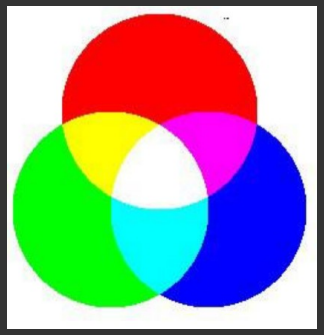

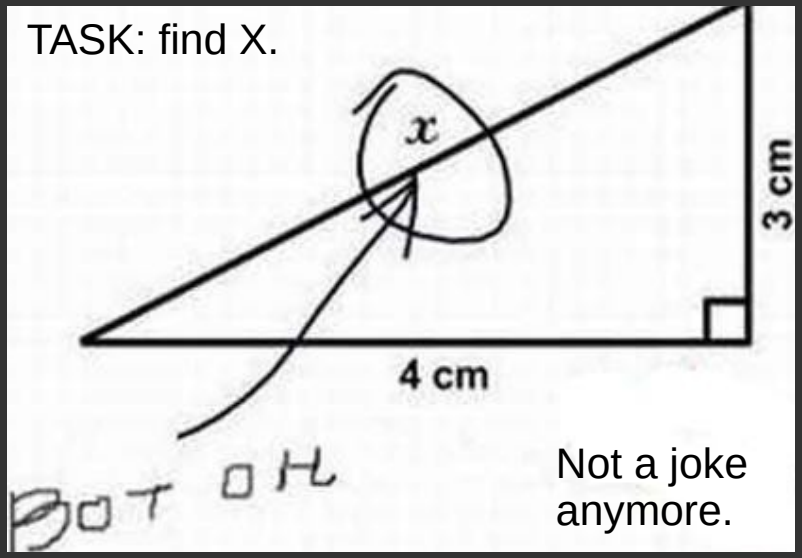

General view of the neural network

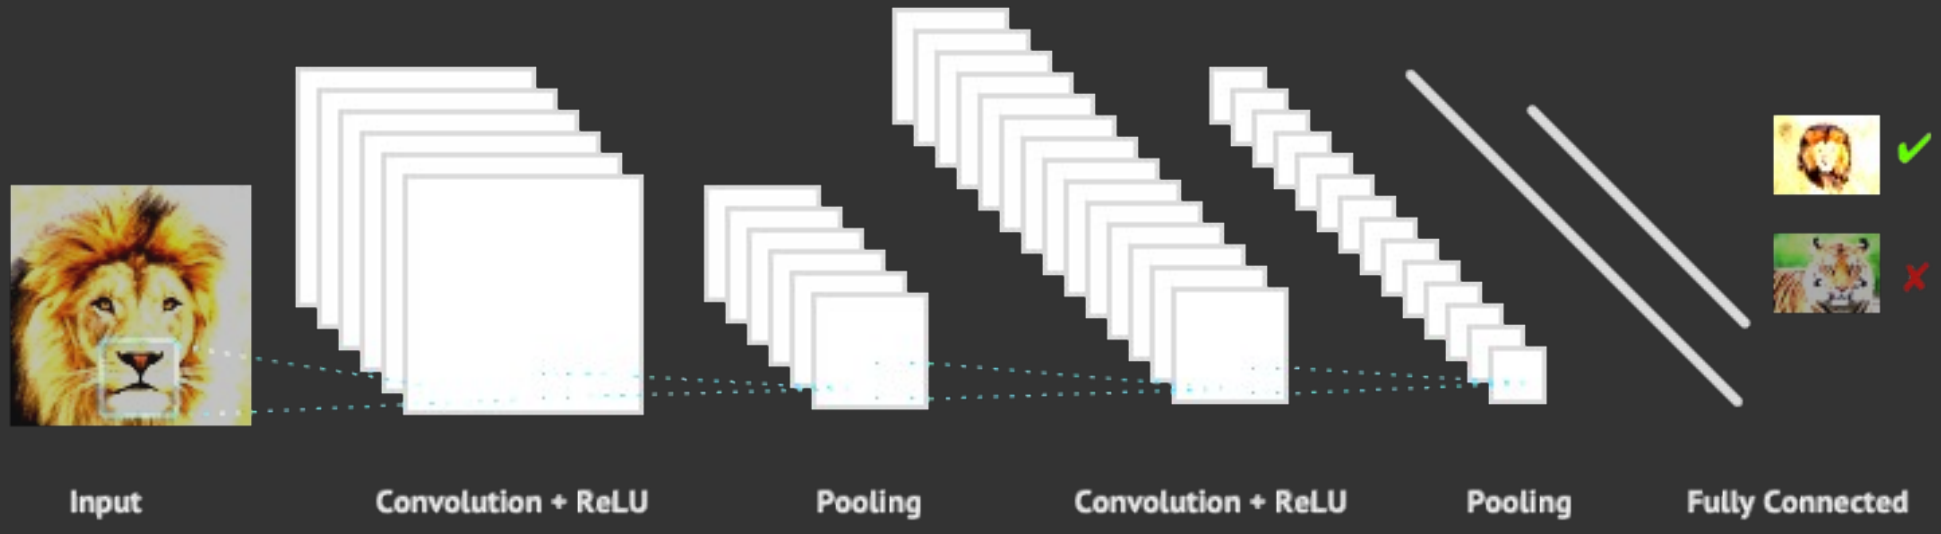

General view of the neural network

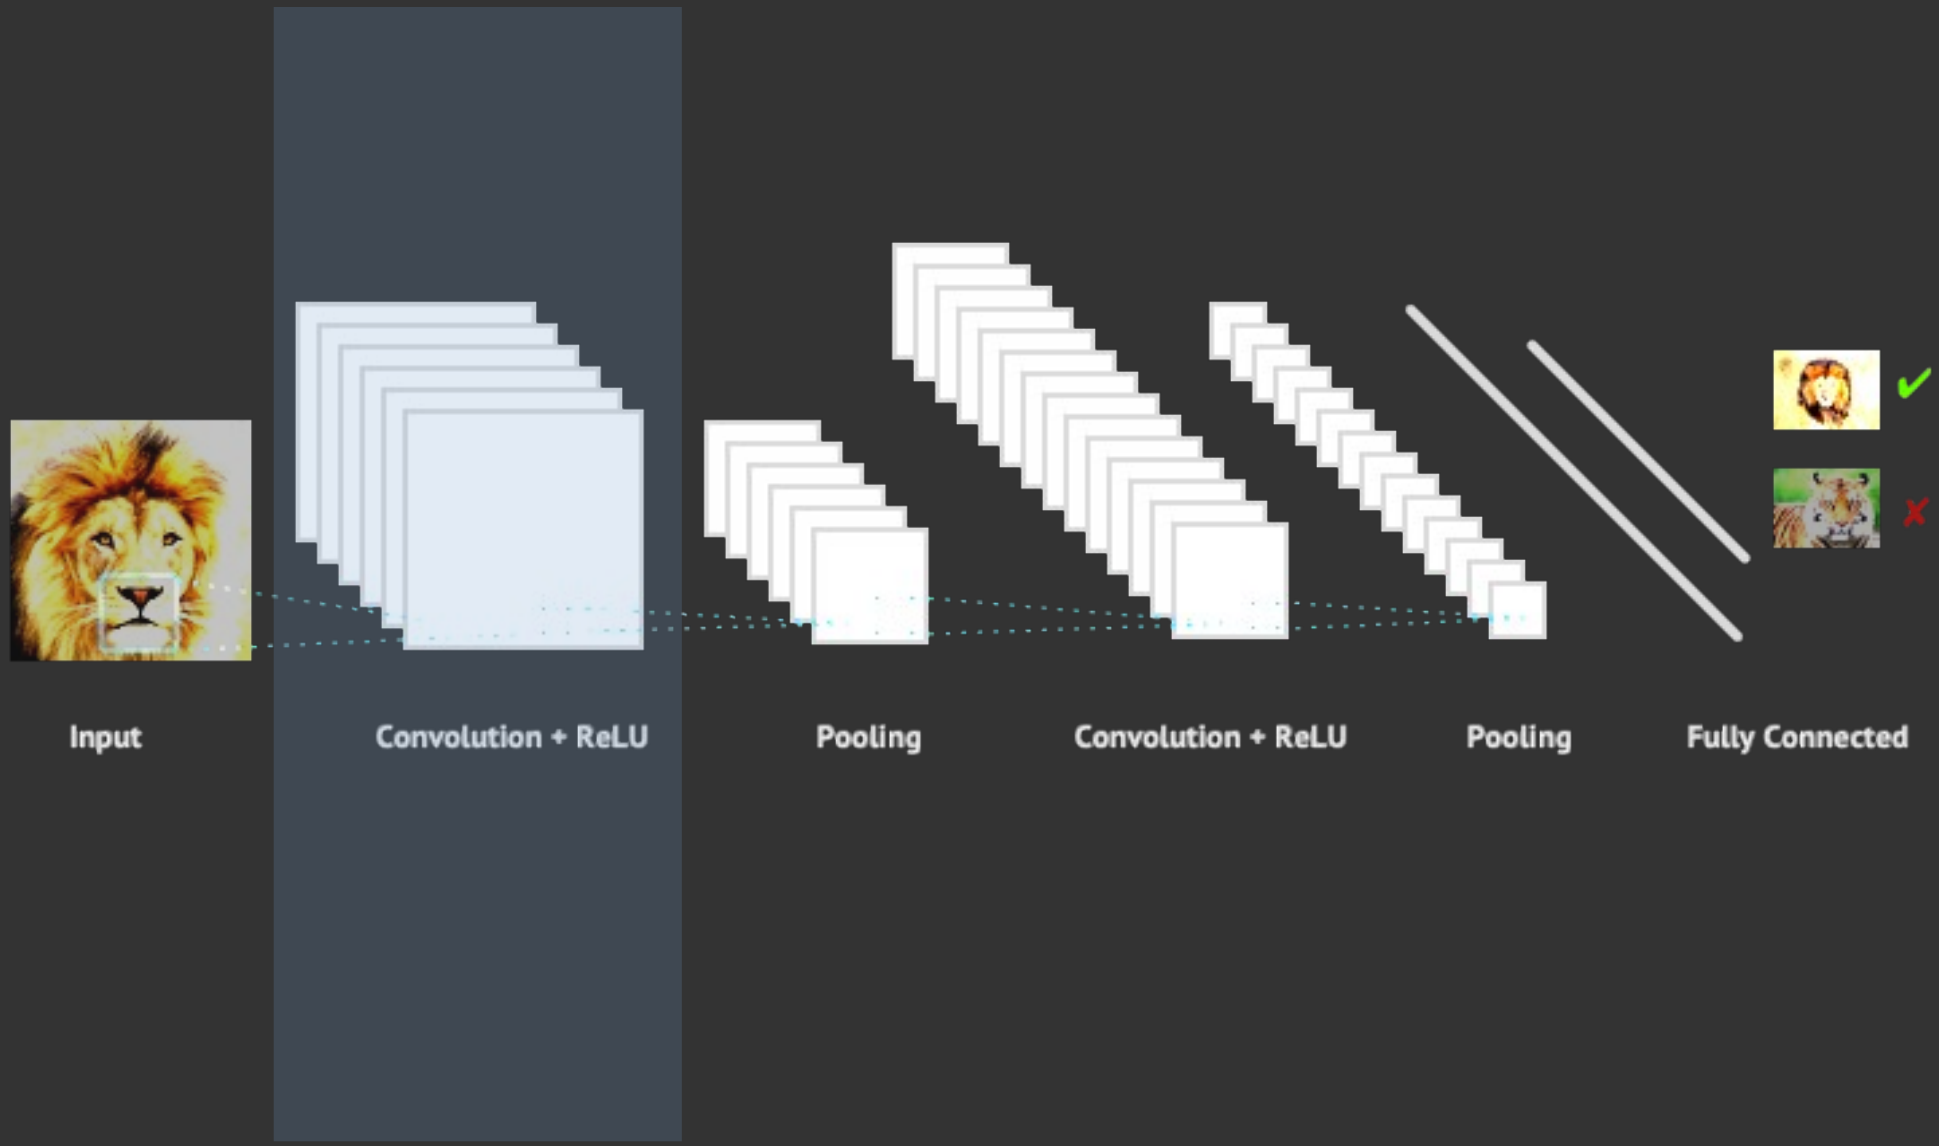

Convolution

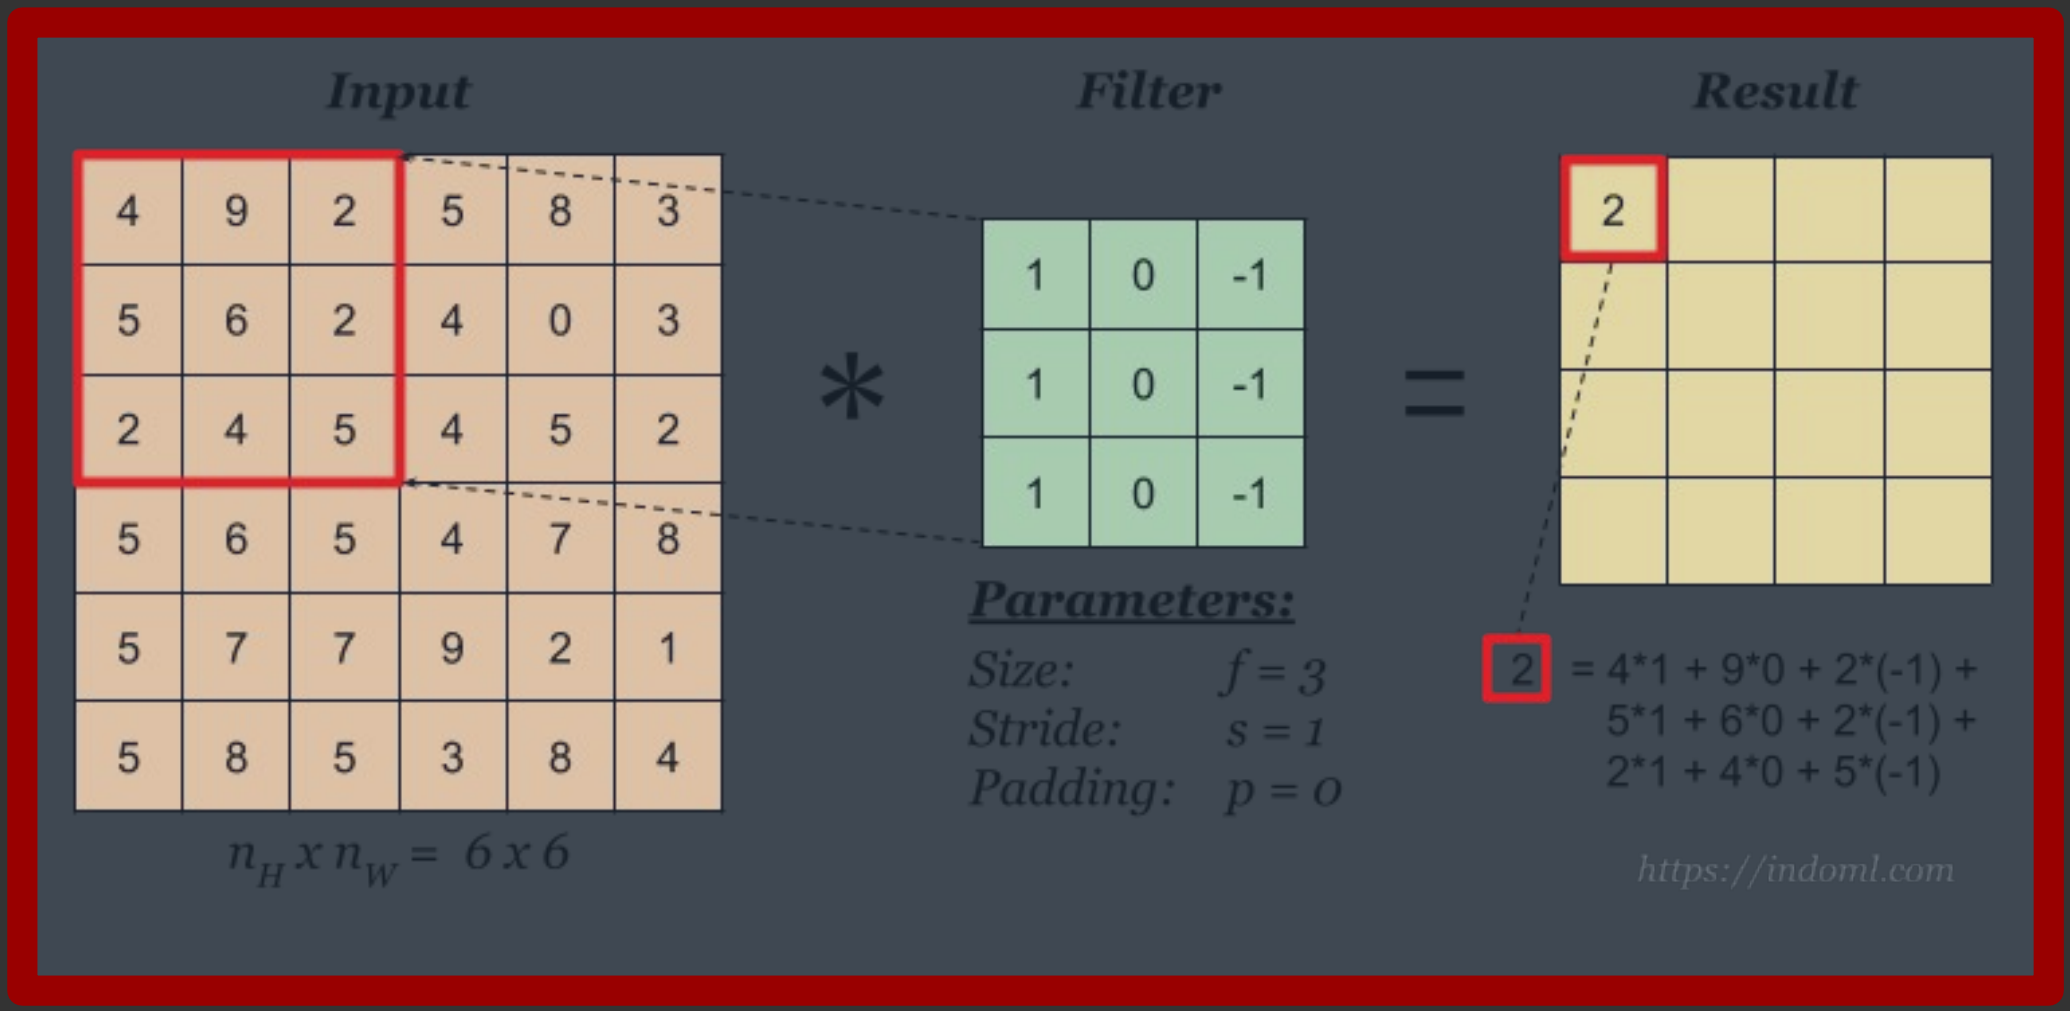

Kernel, or convolution matrix, or mask

Operation	Kernel ω	Image result g(x,y)
Identity	$\begin{bmatrix} 0 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 0 \end{bmatrix}$	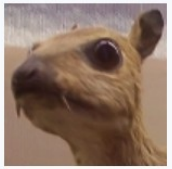
Edge detection	$\begin{bmatrix} 1 & 0 & -1 \\ 0 & 0 & 0 \\ -1 & 0 & 1 \end{bmatrix}$	
	$\begin{bmatrix} 0 & -1 & 0 \\ -1 & 4 & -1 \\ 0 & -1 & 0 \end{bmatrix}$	
	$\begin{bmatrix} -1 & -1 & -1 \\ -1 & 8 & -1 \\ -1 & -1 & -1 \end{bmatrix}$	
Sharpen	$\begin{bmatrix} 0 & -1 & 0 \\ -1 & 5 & -1 \\ 0 & -1 & 0 \end{bmatrix}$	

General view of the neural network

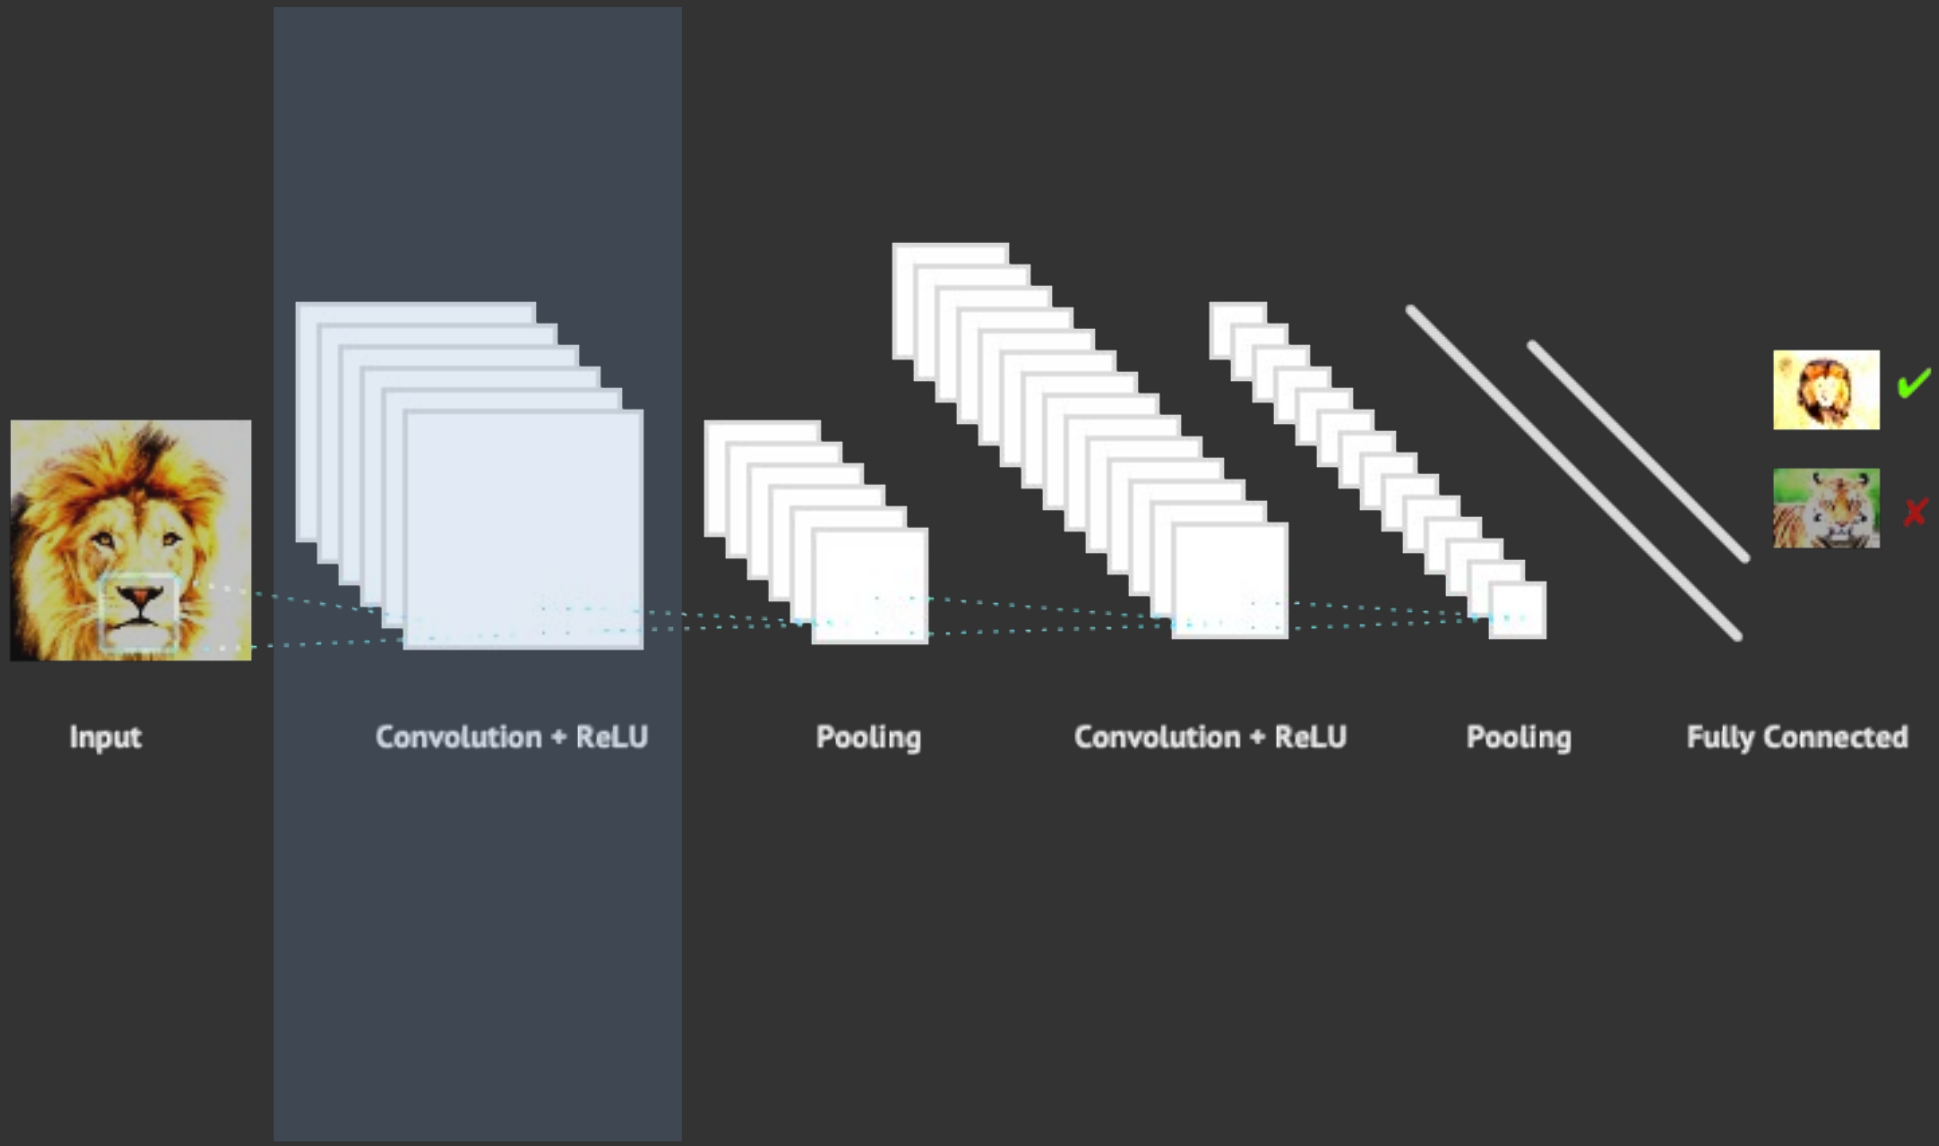

Rectified Linear Unit (ReLU)
(activation function; optimization of calculations)

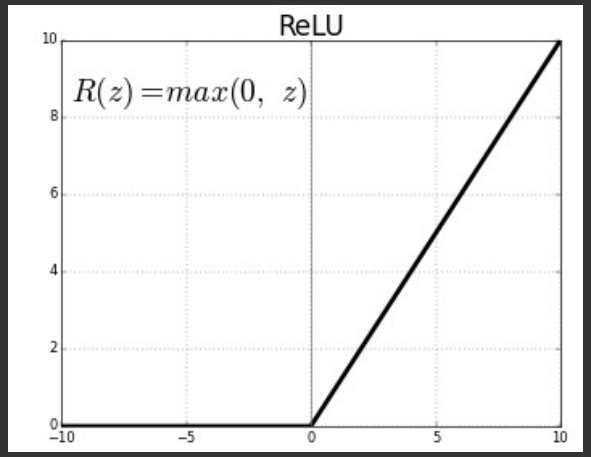

0.77	-0.11	0.11	0.33	0.55	-0.11	0.33
-0.11	1.00	-0.11	0.33	-0.11	0.11	-0.11
0.11	-0.11	1.00	-0.33	0.11	-0.11	0.55
0.33	0.33	-0.33	0.55	-0.33	0.33	0.33
0.55	-0.11	0.11	-0.33	1.00	-0.11	0.11
-0.11	0.11	-0.11	0.33	-0.11	1.00	-0.11
0.33	-0.11	0.55	0.33	0.11	-0.11	0.77

0.77	0	0.11	0.33	0.55	0	0.33
0	1.00	0	0.33	0	0.11	0
0.11	0	1.00	0	0.11	0	0.55
0.33	0.33	0	0.55	0	0.33	0.33
0.55	0	0.11	0	1.00	0	0.11
0	0.11	0	0.33	0	1.00	0
0.33	0	0.55	0.33	0.11	0	0.77

General view of the neural network

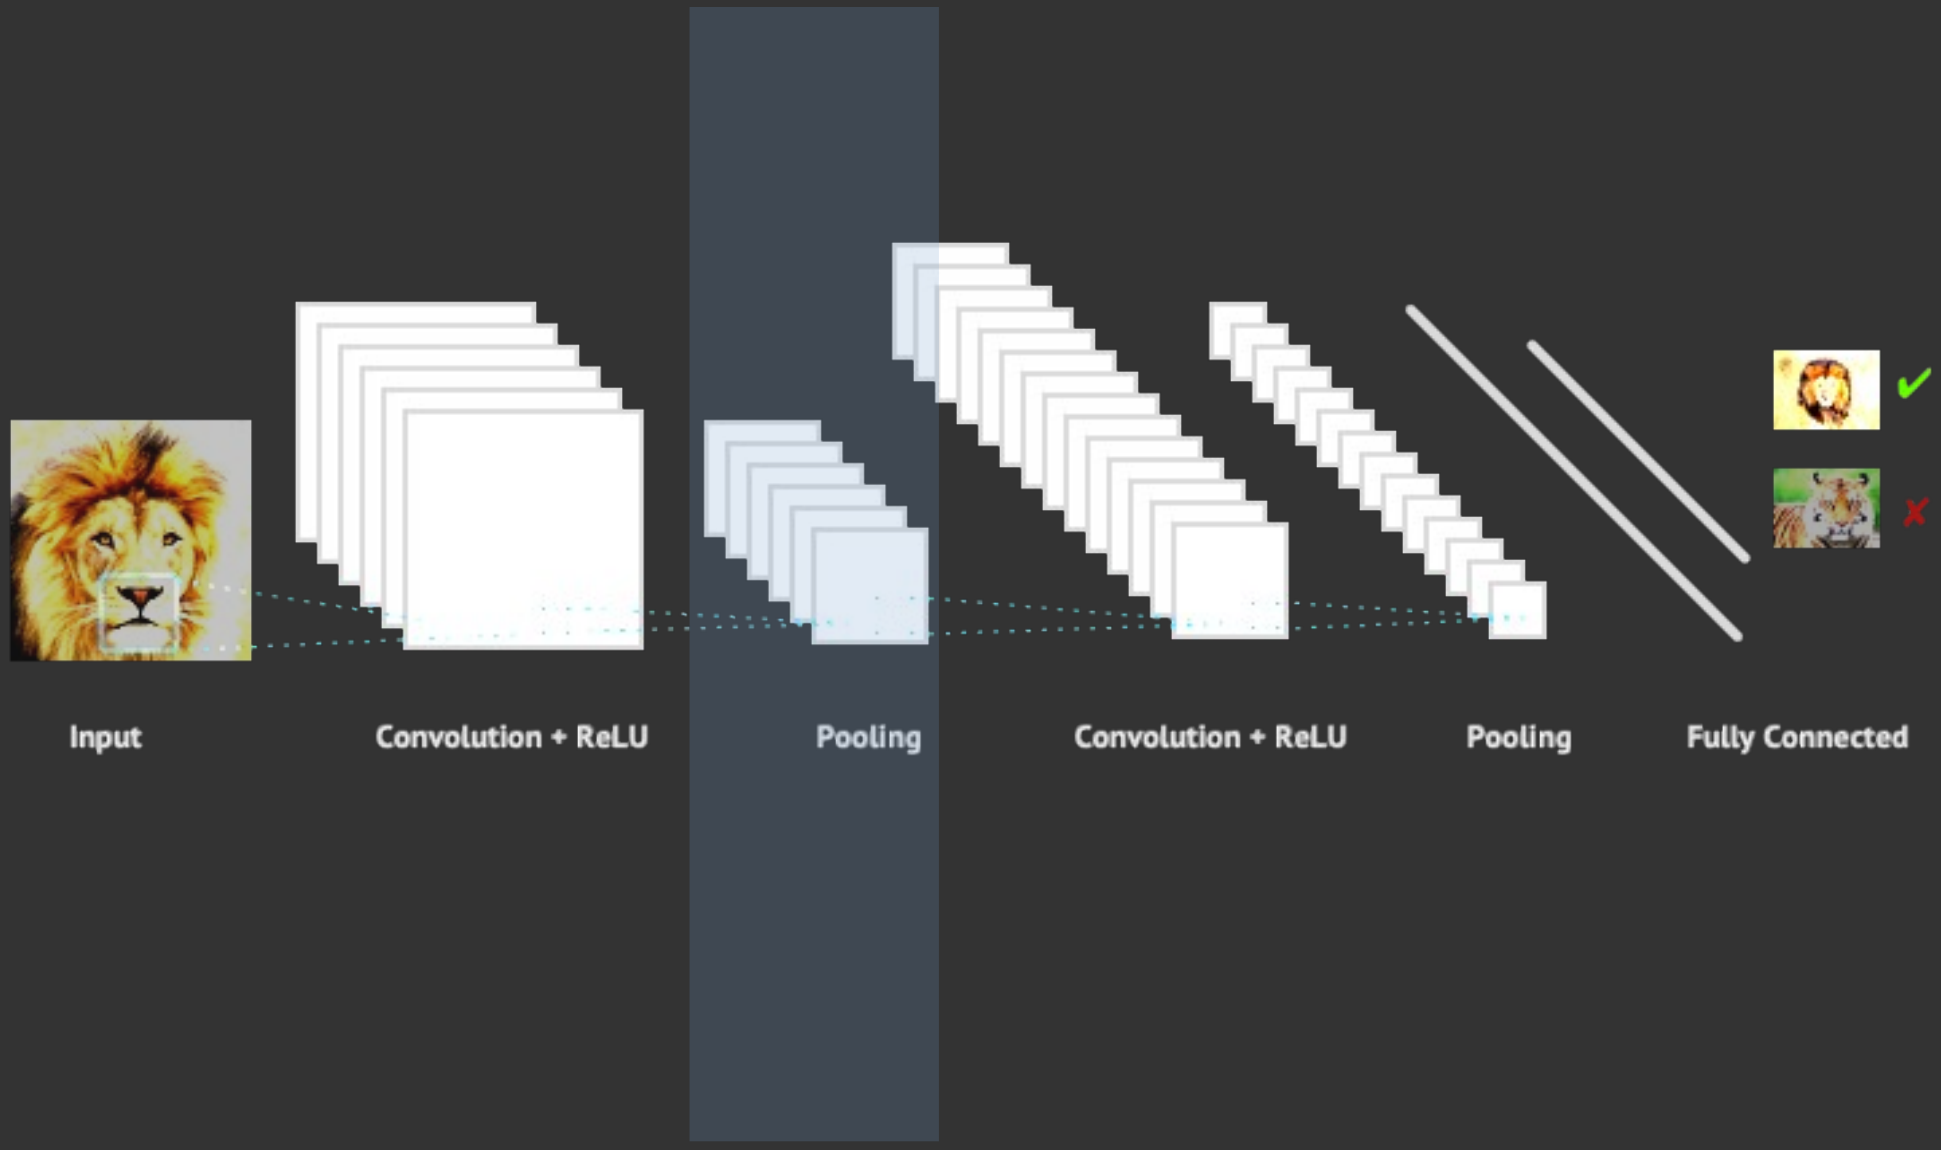

Pooling
(dimension reduction)

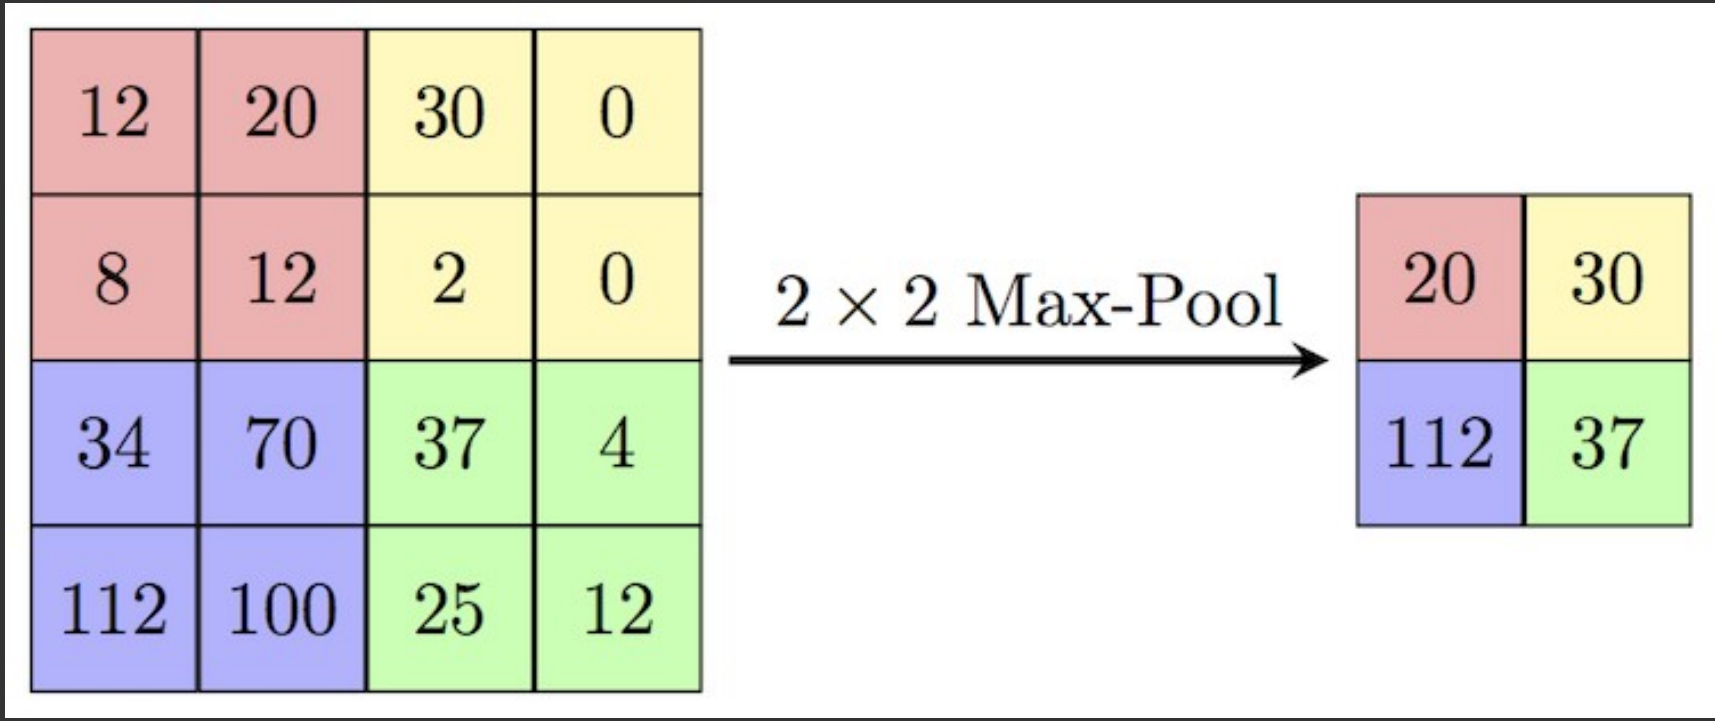

General view of the neural network

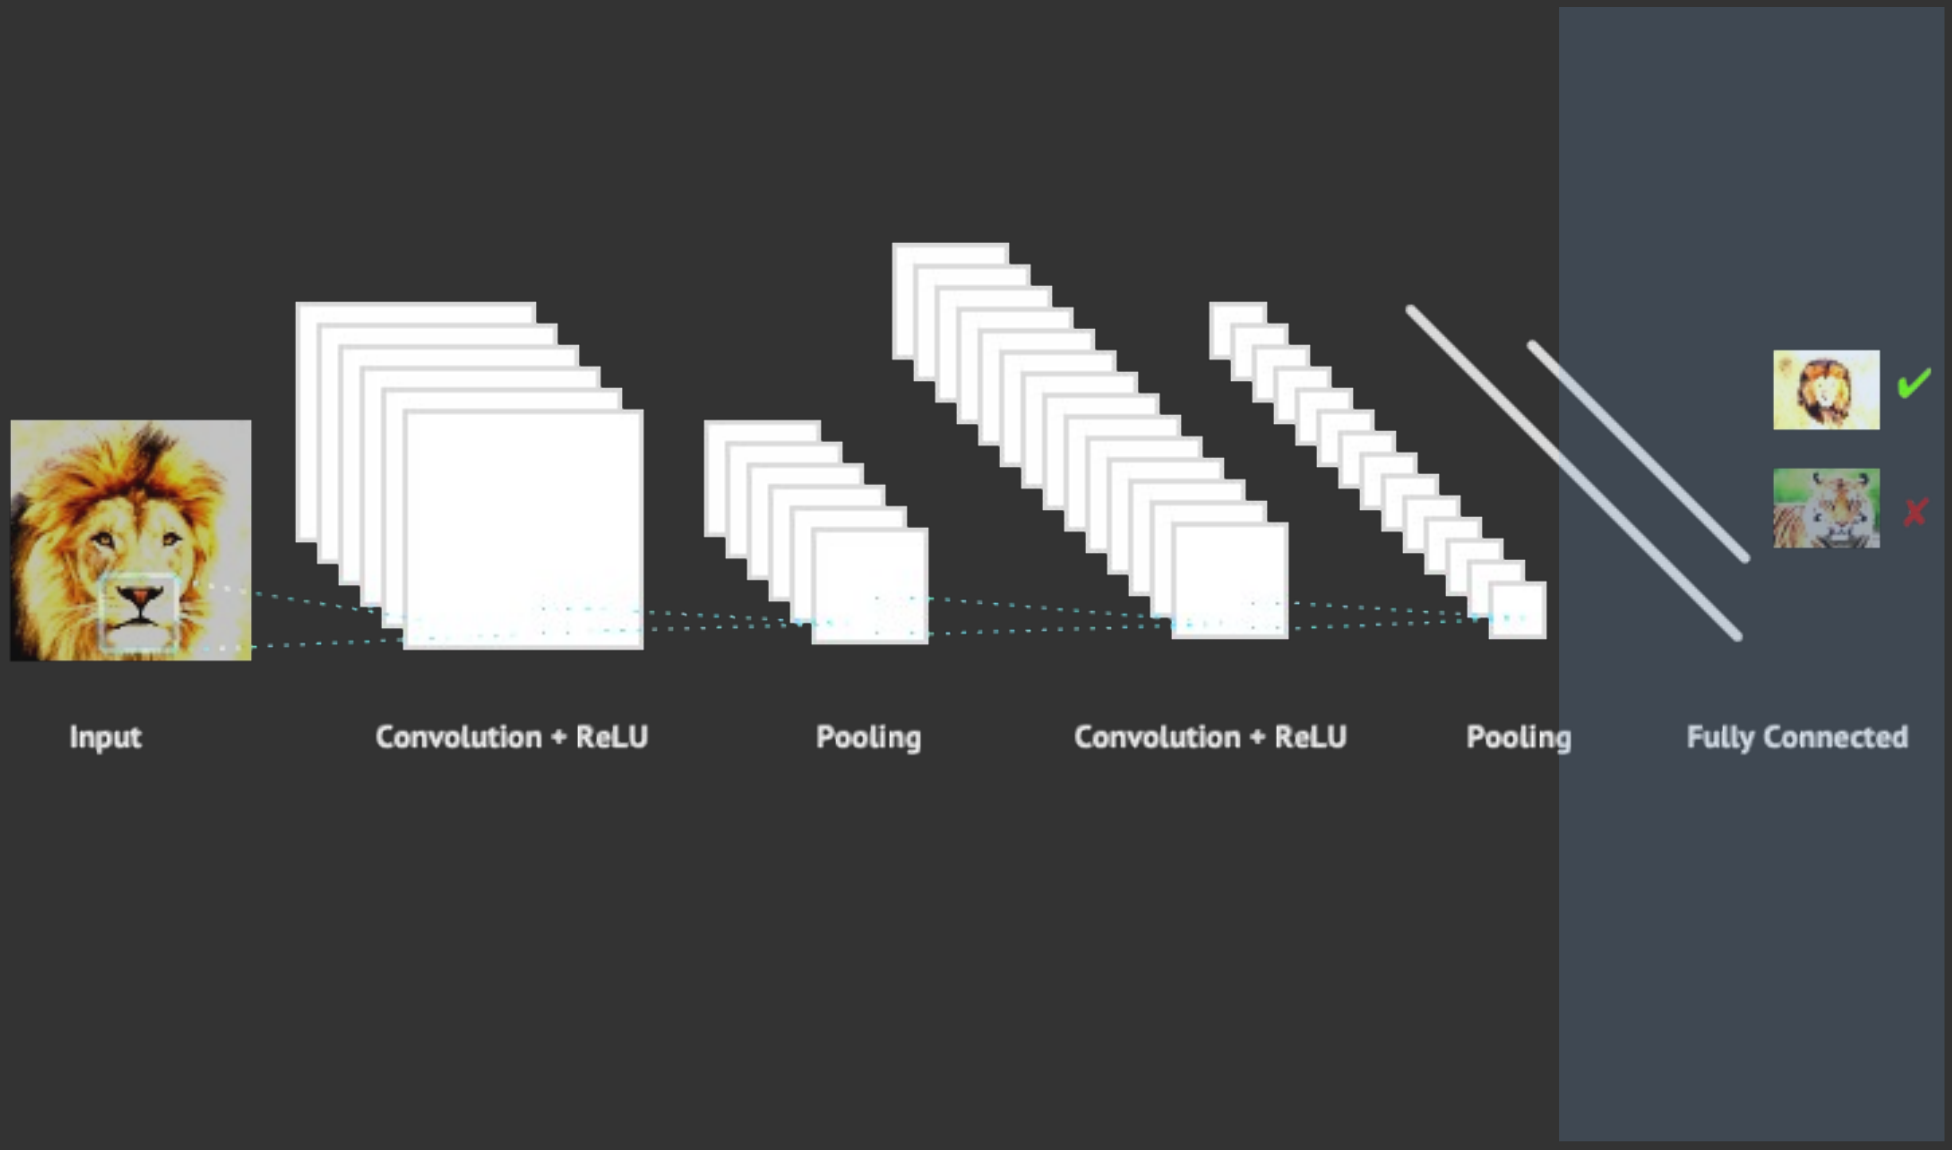

Fully connected (dense) layer

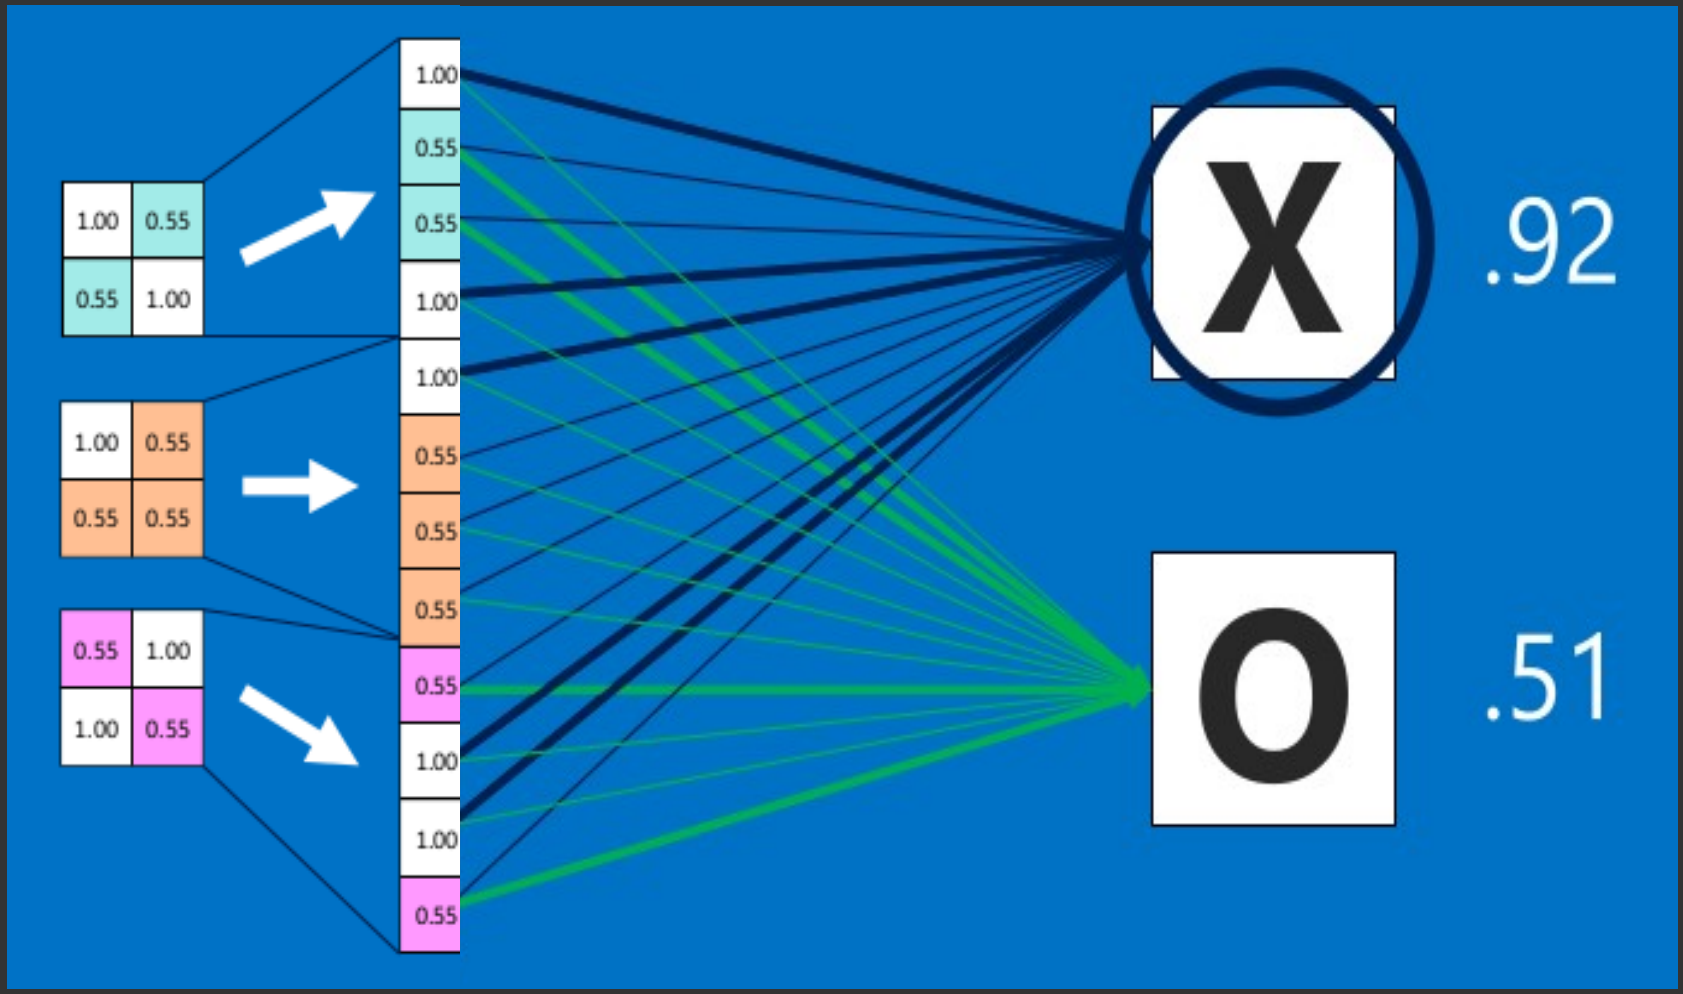

Dropout
(regularization prevents overfitting)

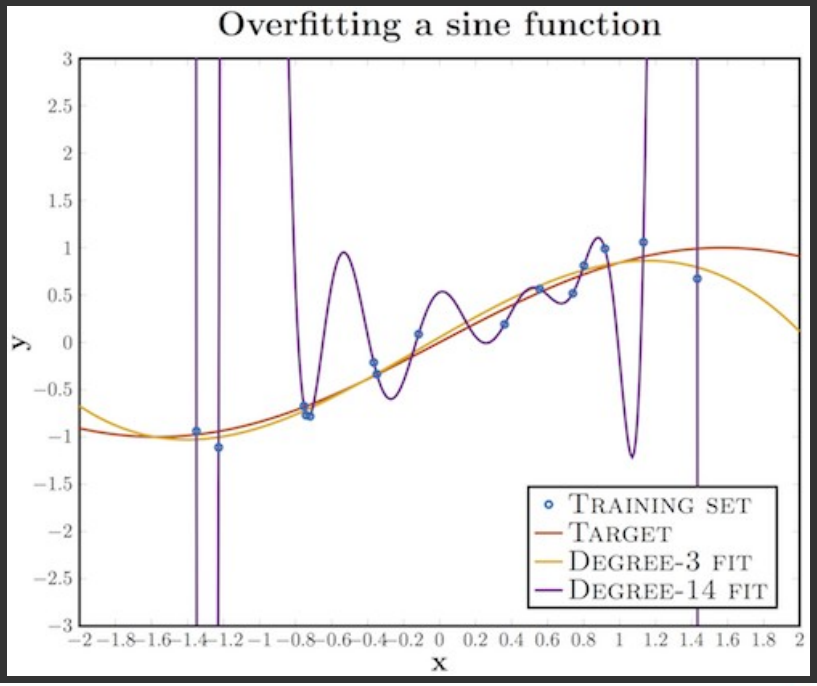

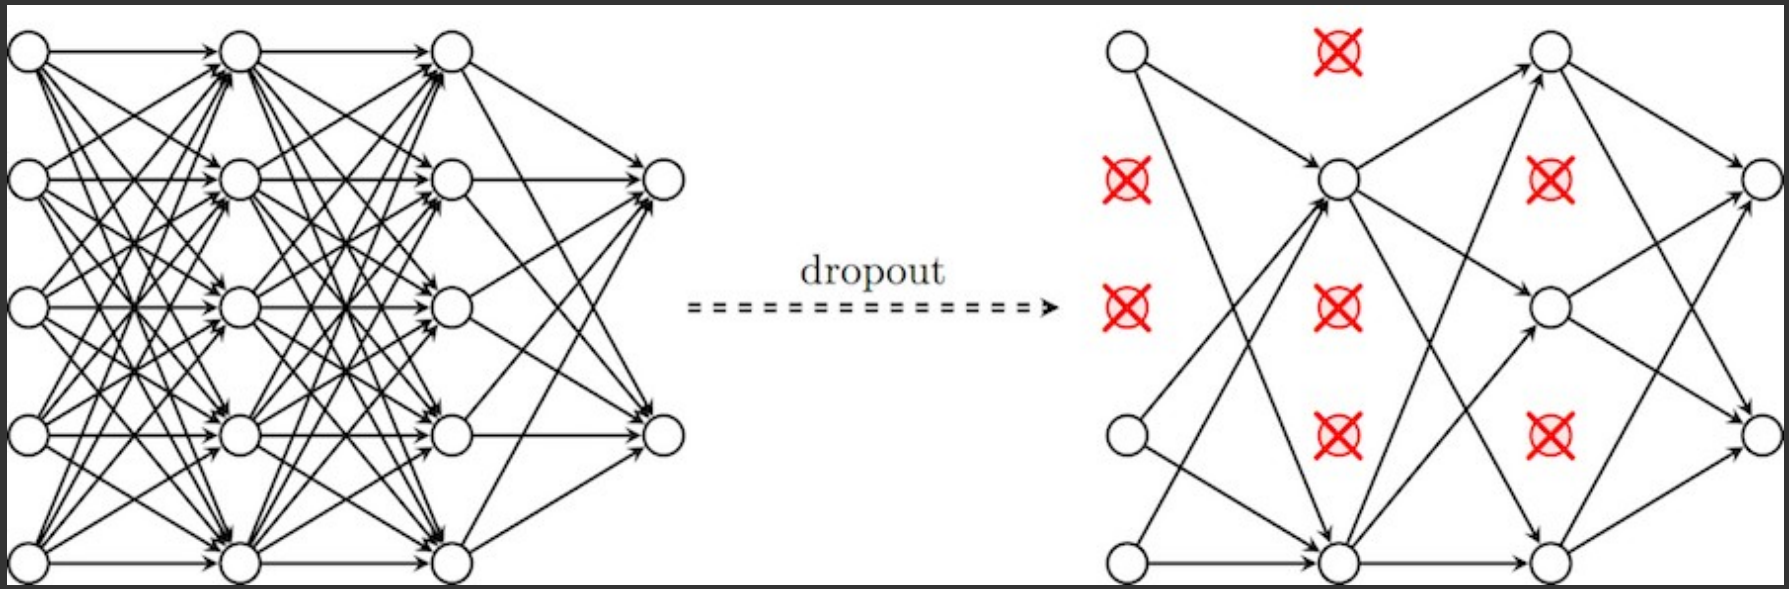

Playing with neural networks (tuning hyperparameters, etc)

http://playground.tensorflow.org

Backpropagation and learning rate